

# CV overview and R-CNN

Vladislav Goncharenko  
Spring 2019  
MIPT, Moscow

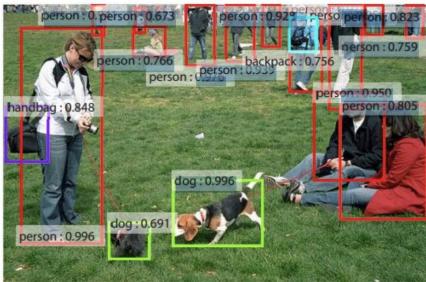
# Outline

- Computer Vision tasks
  - Classification
  - Localisation
  - Detection
  - Semantic Segmentation
  - Instance Segmentation
- Metrics
- Datasets
- R-CNN
- Fast R-CNN
- Faster R-CNN

# Computer Vision tasks



# Computer Vision tasks



- Classification
- Localisation
- Detection
- Semantic Segmentation
- Instance Segmentation
- Human Pose Estimation
- Person Reidentification
- Object Tracking
- Style Transfer
- Adversarial Attacks
- .....

# Classification



Cat

**Input** image ( $C \times H \times W$ )

**Output**

class label (int)

# Localisation



Cat

**Input** image ( $C \times H \times W$ )

**Output**

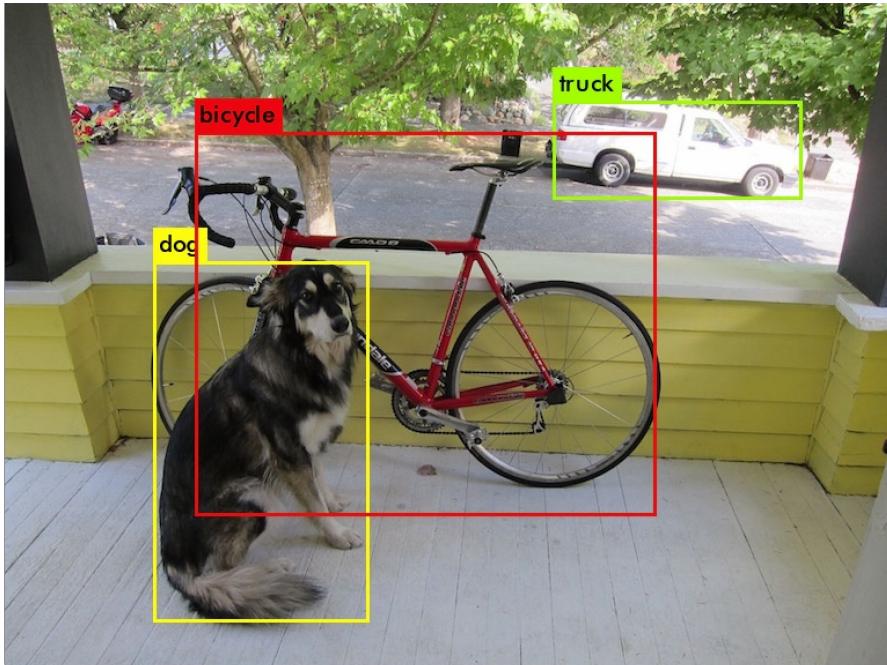
class label (int), BB\* coordinates (4 ints)

**Assumption**

one object per image

\* BB - bounding box

# Detection



**Input** image ( $C \times H \times W$ )

**Output**

class label (int), confidence (0-1 float), BB (4 ints)

.....

class label (int), confidence (0-1 float), BB (4 ints)

**Assumption**

multiple objects per image

# Semantic Segmentation



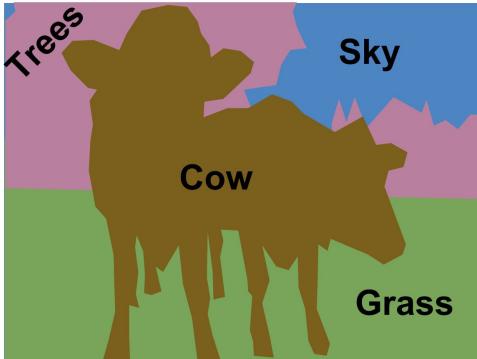
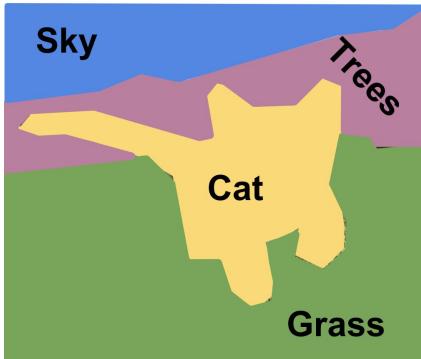
Input image ( $C \times H \times W$ )

Output

mask of classes ( $1 \times H \times W$ )

Assumption

just class label for each pixel



# Instance Segmentation



**Input** image ( $C \times H \times W$ )

**Output**

mask of class instances (detection + 1  $\times H \times W$ )

**Assumption**

separated masks for each distinct object

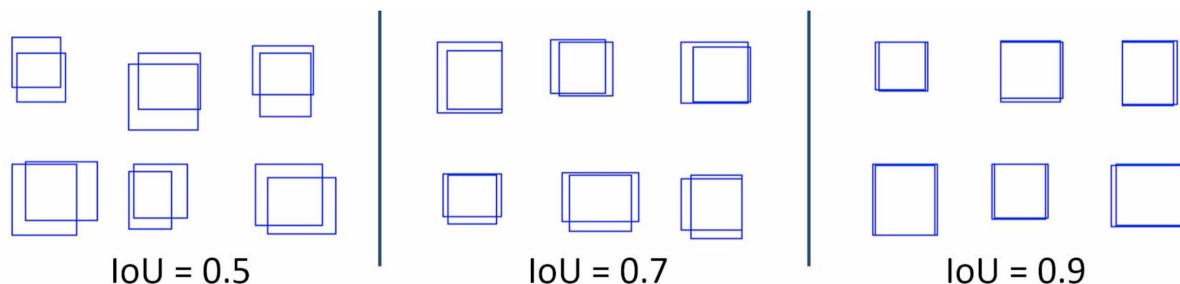
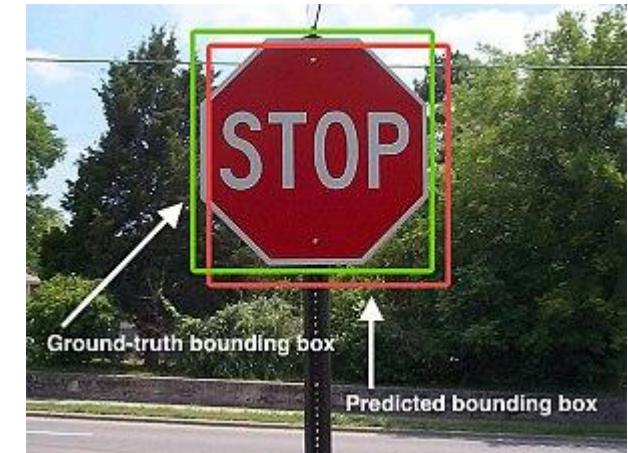
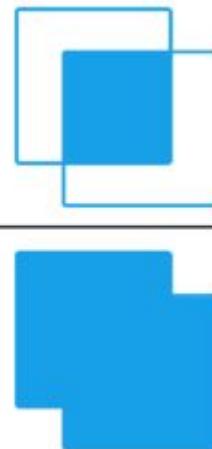
# Metrics in CV



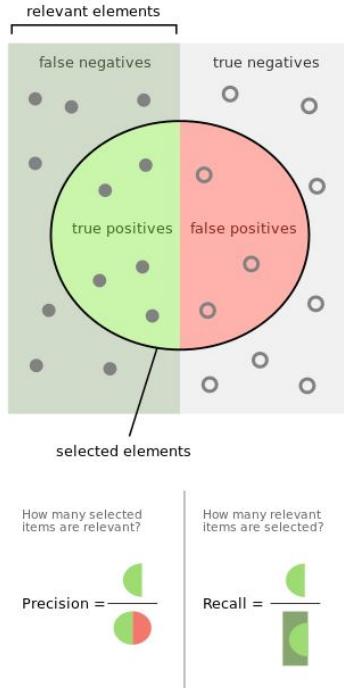
# Intersection over Union (IoU)

aka Jaccard index

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$



# mean Average Precision (mAP)

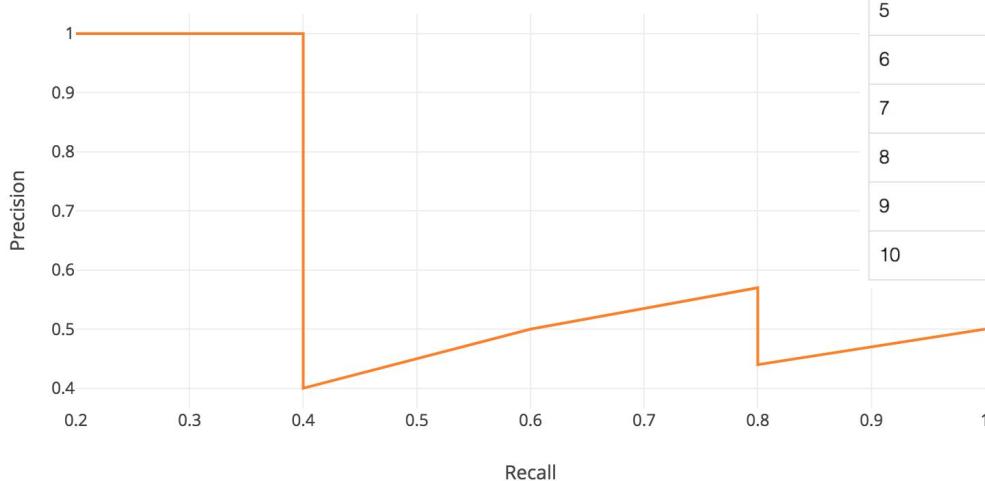


Consider 5 objects and 10 detections

Rank	Correct?	Precision	Recall
1	True	1.0	0.2
2	True	1.0	0.4
3	False	0.67	0.4
4	False	0.5	0.4
5	False	0.4	0.4
6	True	0.5	0.6
7	True	0.57	0.8
8	False	0.5	0.8
9	False	0.44	0.8
10	True	0.5	1.0

# mean Average Precision (mAP)

The intention in interpolating the precision/recall curve in this way is to reduce the impact of the “wiggles” in the precision/recall curve, caused by small variations in the ranking of examples.

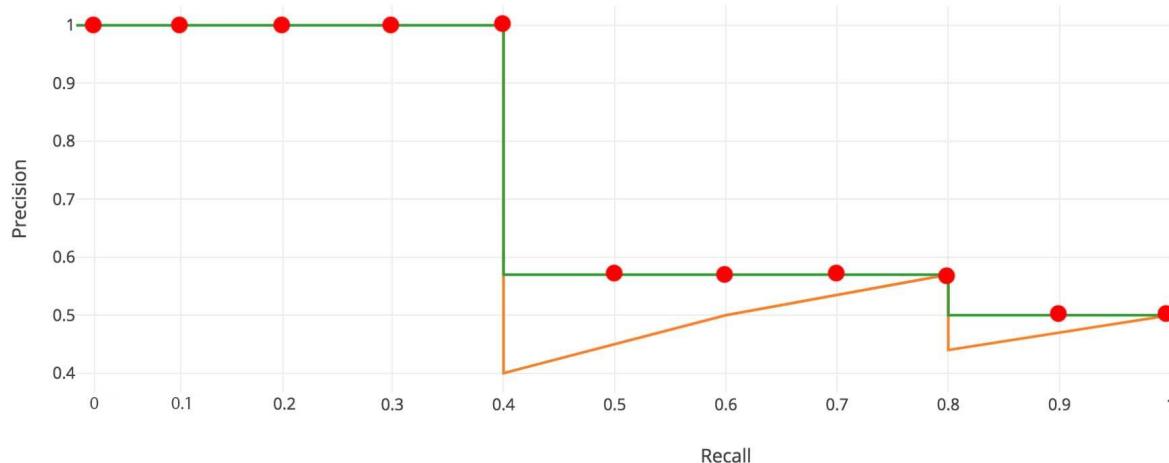


Rank	Correct?	Precision	Recall
1	True	1.0	0.2
2	True	1.0	0.4
3	False	0.67	0.4
4	False	0.5	0.4
5	False	0.4	0.4
6	True	0.5	0.6
7	True	0.57	0.8
8	False	0.5	0.8
9	False	0.44	0.8
10	True	0.5	1.0

# mean Average Precision (mAP)

mAP is mean of AP over all classes

$$AP = \frac{1}{11} \times (AP_r(0) + AP_r(0.1) + \dots + AP_r(1.0))$$

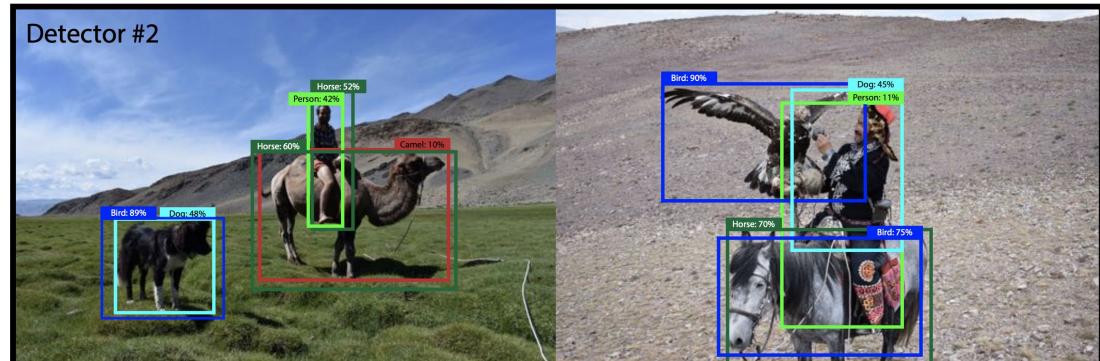
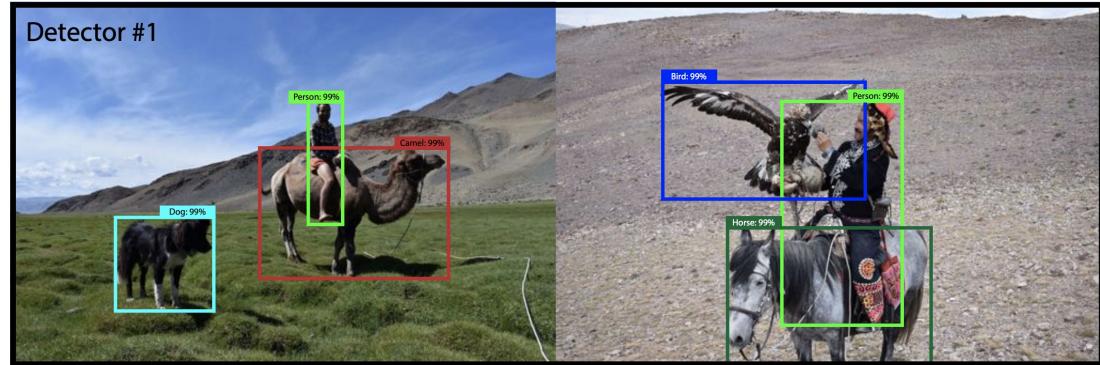


# mean Average Precision (mAP)

## Problems

These two hypothetical detectors are perfect according to mAP over these two images.  
They are both perfect. Totally equal

([from YOLOv3 article](#))

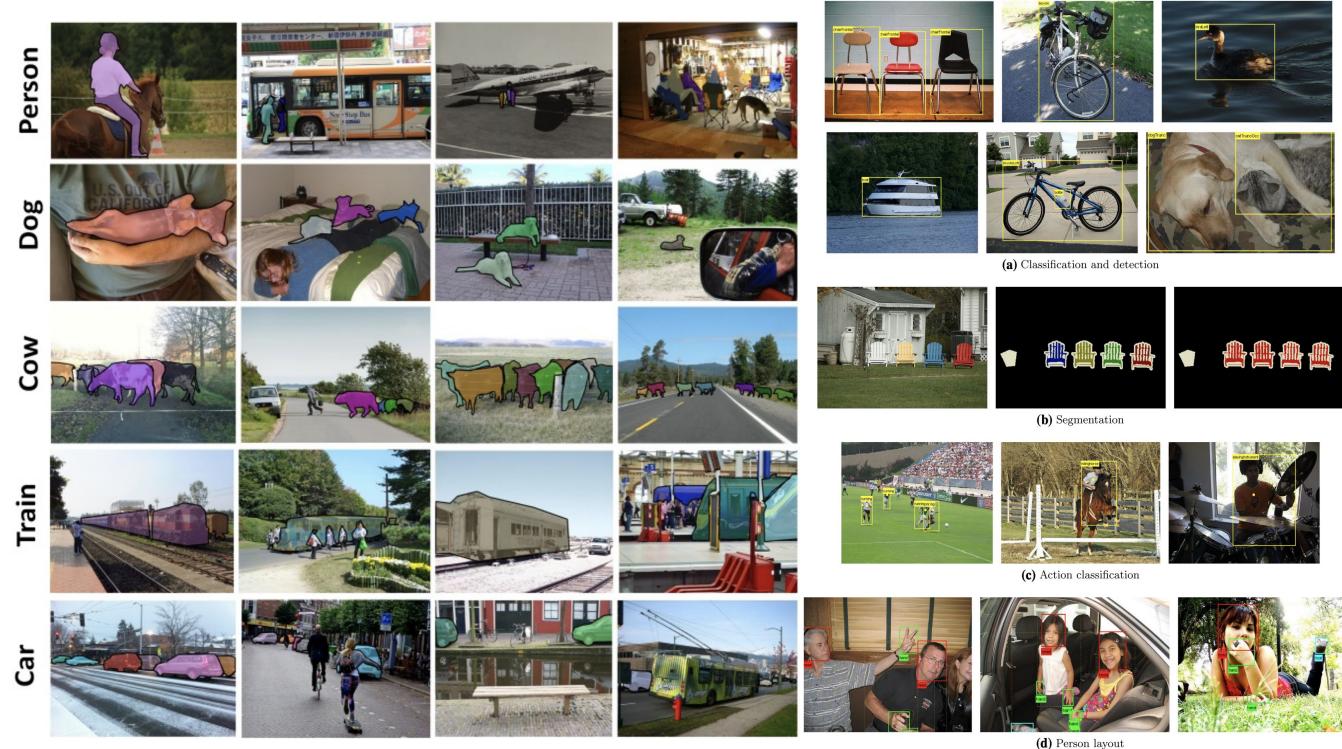


# Datasets



# Detection datasets

1. PASCAL VOC 2012
2. ImageNet
3. MS COCO 2017
4. Google Open Images v4



Thanks to Ilya Zakharkin  
for materials

# PASCAL VOC 2012



**Table 1** The VOC classes

Vehicles	Household	Animals	Other
Aeroplane	Bottle	Bird	Person
Bicycle	Chair	Cat	
Boat	Dining table	Cow	
Bus	Potted plant	Dog	
Car	Sofa	Horse	
Motorbike	TV/Monitor	Sheep	
Train			

The classes can be considered in a notional taxonomy

Relevant dataset overview:

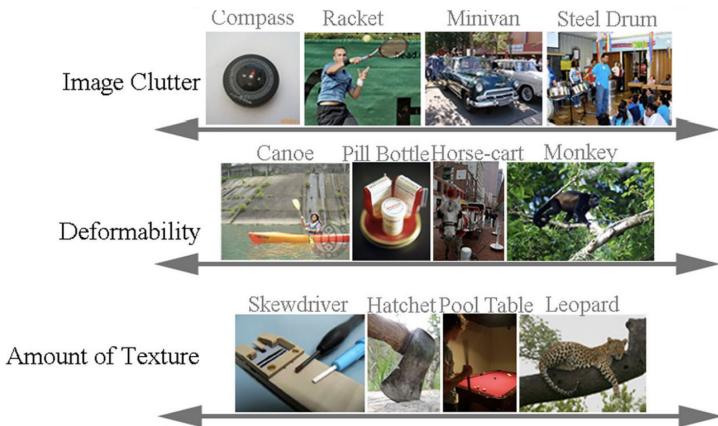
<http://host.robots.ox.ac.uk/pascal/VOC/pubs/everingham15.pdf>

- PASCAL Visual Object Classes
- developed since 2005 and 4 classes
- 11,5k images
- 20 classes
- 31,5k detections
- 7k segmentations
- Images came from flickr.com

# ImageNet

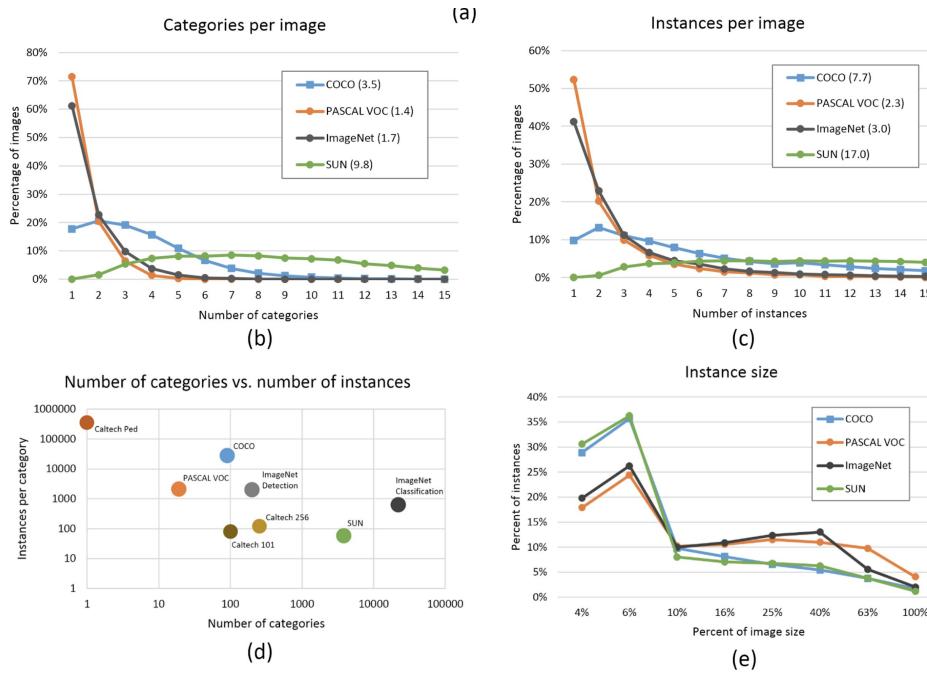


- organized according to the WordNet hierarchy
- 14.2m images
- 1000 classes
- 1m images with BBs
- more fine graded classes than PASCAL
- more diverse objects properties



<http://image-net.org/explore>

# MS COCO 2017



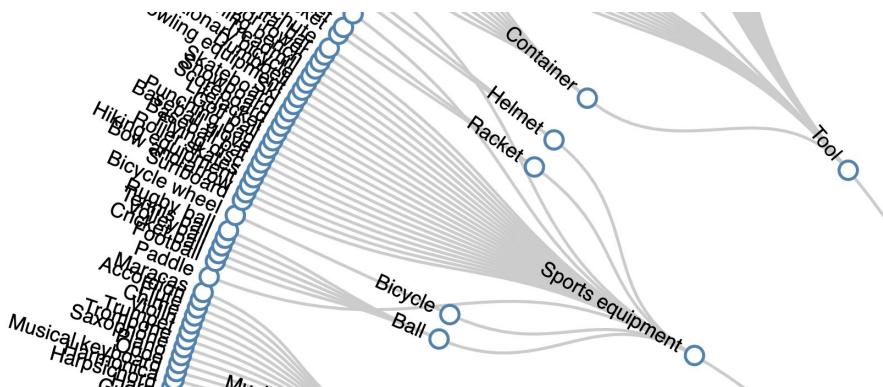
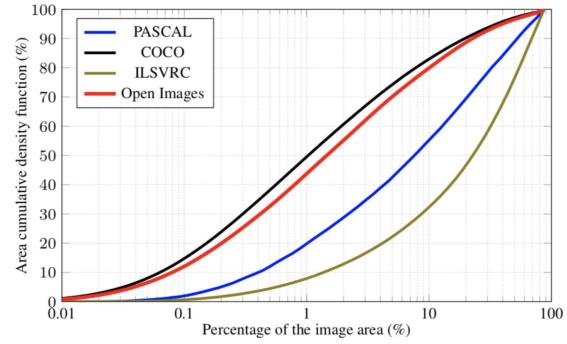
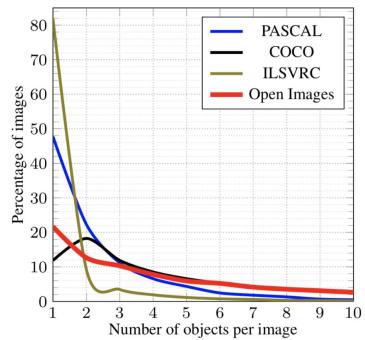
- MicroSoft Common Objects in COnText
- 328k images
- 91 object categories
- 82 having more than 5,000 labeled instances
- 2,5m labeled instances in total
- 250k people with keypoints
- online exploration tool  
<http://cocodataset.org/#explore>
- overview <https://arxiv.org/pdf/1405.0312.pdf>

# Open Images

- Currently relevant challenge
- 15m boxes on 600 categories
- 2.8m instance segmentations on 350 categories
- 36.5m image-level labels on 20k categories
- 391k relationship annotations of 329 relationships
- Extension - 478k crowdsourced images with 6k+ categories



<https://storage.googleapis.com/openimages/web/factsfigures.html>



# R-CNN



# Detection as Classification



**CAT? NO**

**DOG? NO**

# Detection as Classification

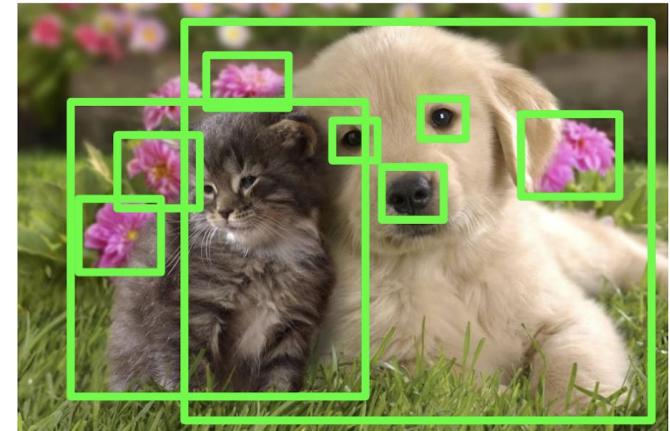


CAT? YES!

DOG? NO

# Region Proposals

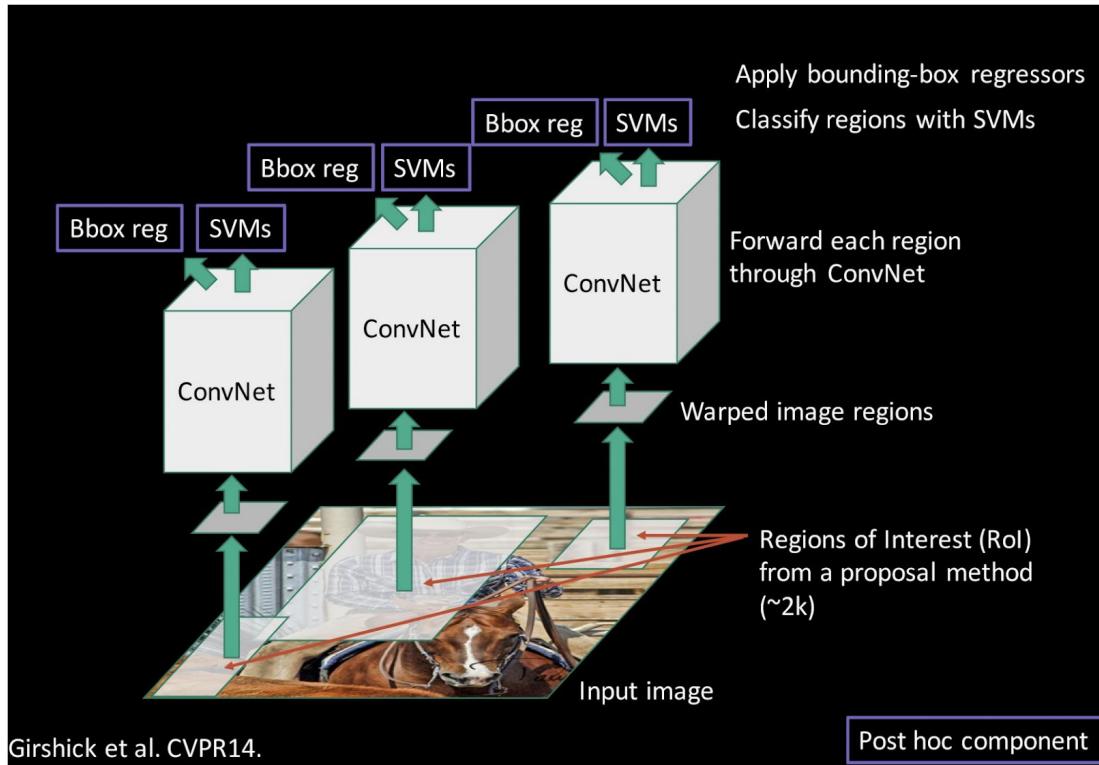
- Find “blobby” image regions that are likely to contain objects
- “Class-agnostic” object detector
- Look for “blob-like” regions



# Region Proposals: Many other choices

Method	Approach	Outputs Segments	Outputs Score	Control #proposals	Time (sec.)	Repea- tability	Recall Results	Detection Results
Bing [18]	Window scoring		✓	✓	0.2	***	*	.
CPMC [19]	Grouping	✓	✓	✓	250	-	**	*
EdgeBoxes [20]	Window scoring		✓	✓	0.3	**	***	***
Endres [21]	Grouping	✓	✓	✓	100	-	***	**
Geodesic [22]	Grouping	✓		✓	1	*	***	**
MCG [23]	Grouping	✓	✓	✓	30	*	***	***
Objectness [24]	Window scoring		✓	✓	3	.	*	.
Rahtu [25]	Window scoring		✓	✓	3	.	.	*
RandomizedPrim's [26]	Grouping	✓		✓	1	*	*	**
Rantalankila [27]	Grouping	✓		✓	10	**	.	**
Rigor [28]	Grouping	✓		✓	10	*	**	**
SelectiveSearch [29]	Grouping	✓	✓	✓	10	**	***	***
Gaussian				✓	0	.	.	*
SlidingWindow				✓	0	***	.	.
Superpixels		✓			1	*	.	.
Uniform				✓	0	.	.	.

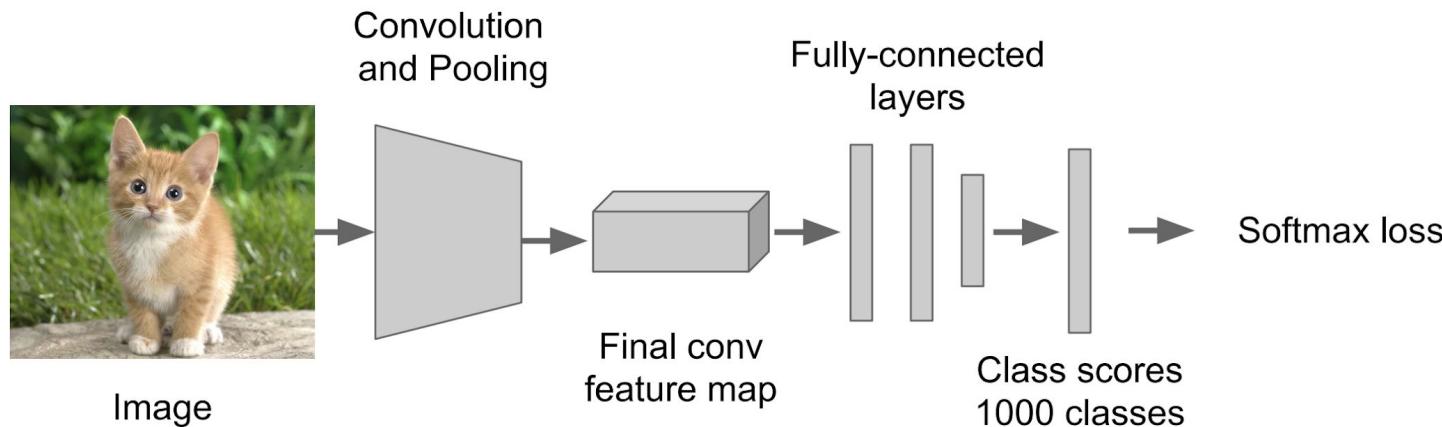
# Putting it together: R-CNN



Girshick et al, “Rich feature hierarchies for accurate object detection and semantic segmentation”, CVPR 2014

# R-CNN Training

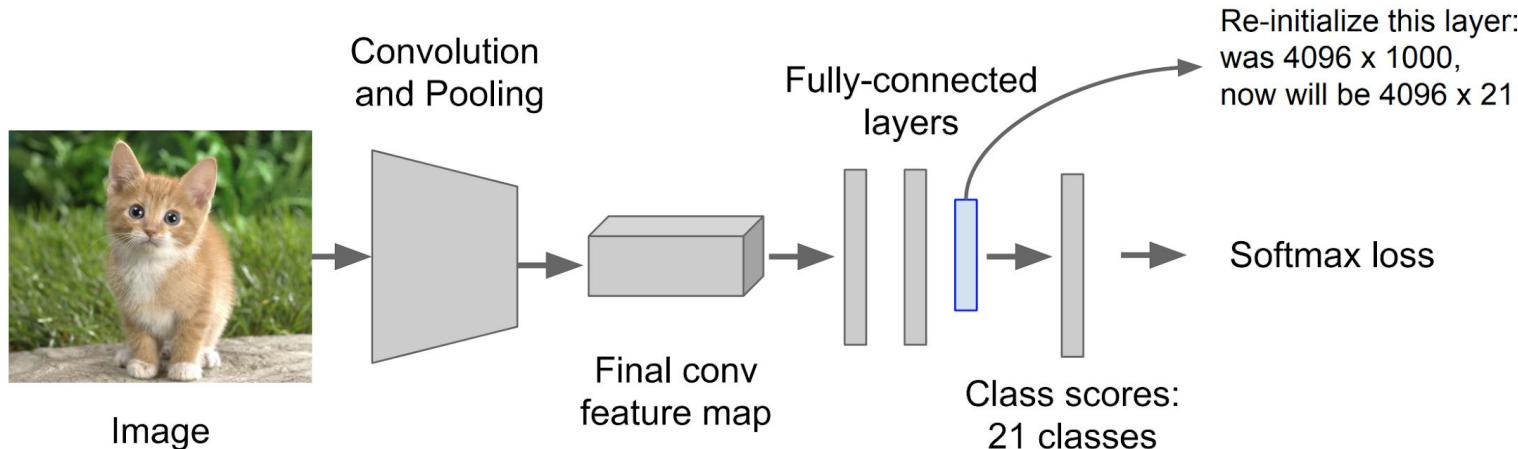
**Step 1:** Train (or download) a classification model for ImageNet (AlexNet)



# R-CNN Training

## Step 2: Fine-tune model for detection

- Instead of 1000 ImageNet classes, want 20 object classes + background
- Throw away final fully-connected layer, reinitialize from scratch
- Keep training model using positive / negative regions from detection images



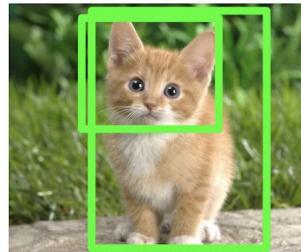
# R-CNN Training

## Step 3: Extract features

- Extract region proposals for all images
- For each region: warp to CNN input size, run forward through CNN, save pool5 features to disk
- Have a big hard drive: features are ~200GB for PASCAL dataset!



Image

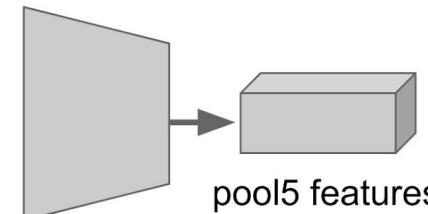


Region Proposals

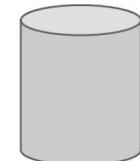


Crop + Warp

Convolution  
and Pooling



pool5 features



Save to disk

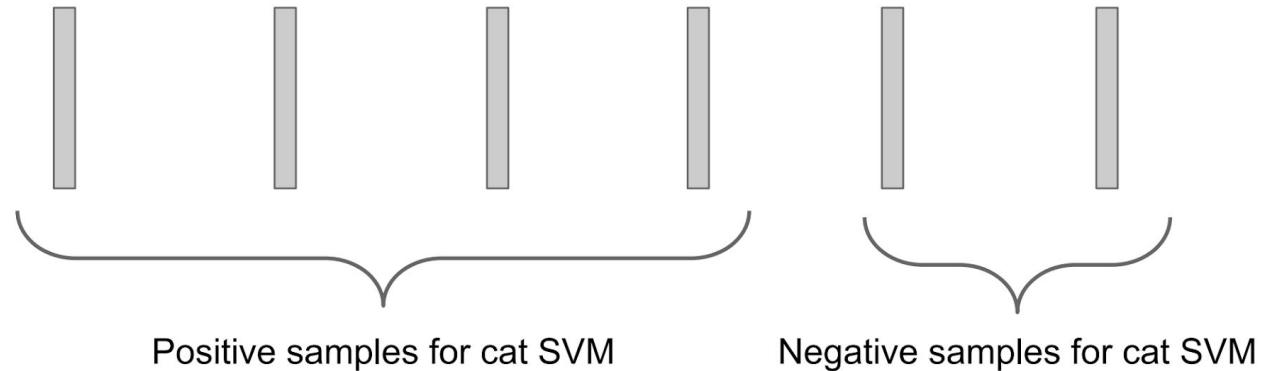
# R-CNN Training

**Step 4:** Train one binary SVM per class to classify region features

Training image regions



Cached region features



# R-CNN Training

**Step 5 (bbox regression):** For each class, train a linear regression model to map from cached features to offsets to GT boxes to make up for “slightly wrong” proposals

Training image regions



Cached region features



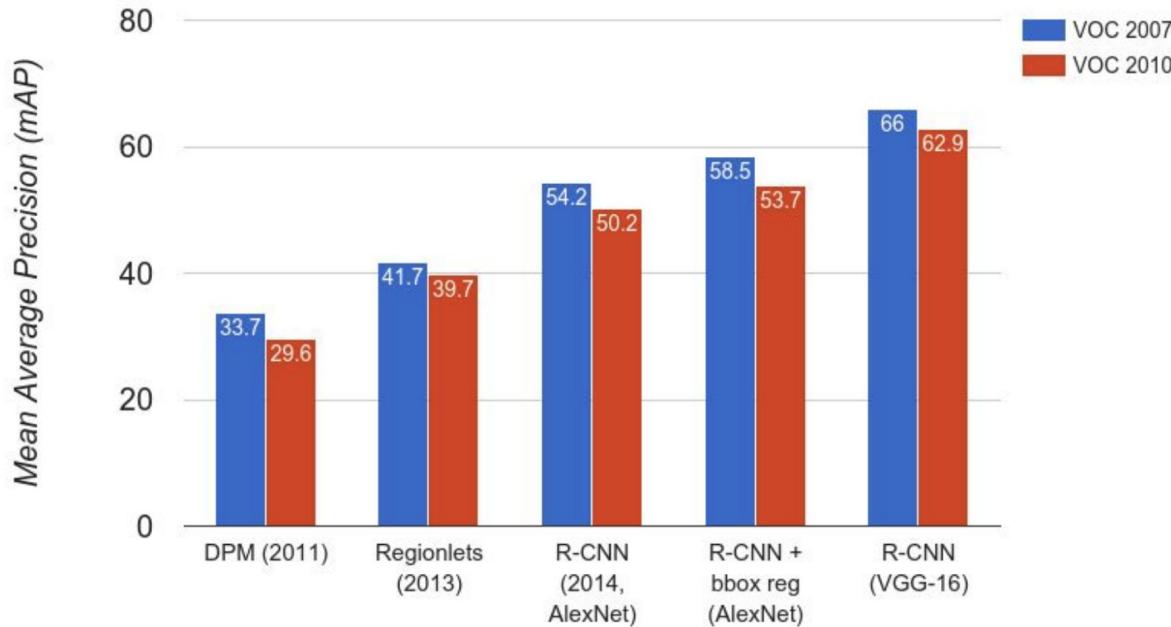
Regression targets  
( $dx$ ,  $dy$ ,  $dw$ ,  $dh$ )  
Normalized coordinates

$(0, 0, 0, 0)$   
Proposal is good

$(.25, 0, 0, 0)$   
Proposal too far to left

$(0, 0, -0.125, 0)$   
Proposal too wide

# R-CNN Results

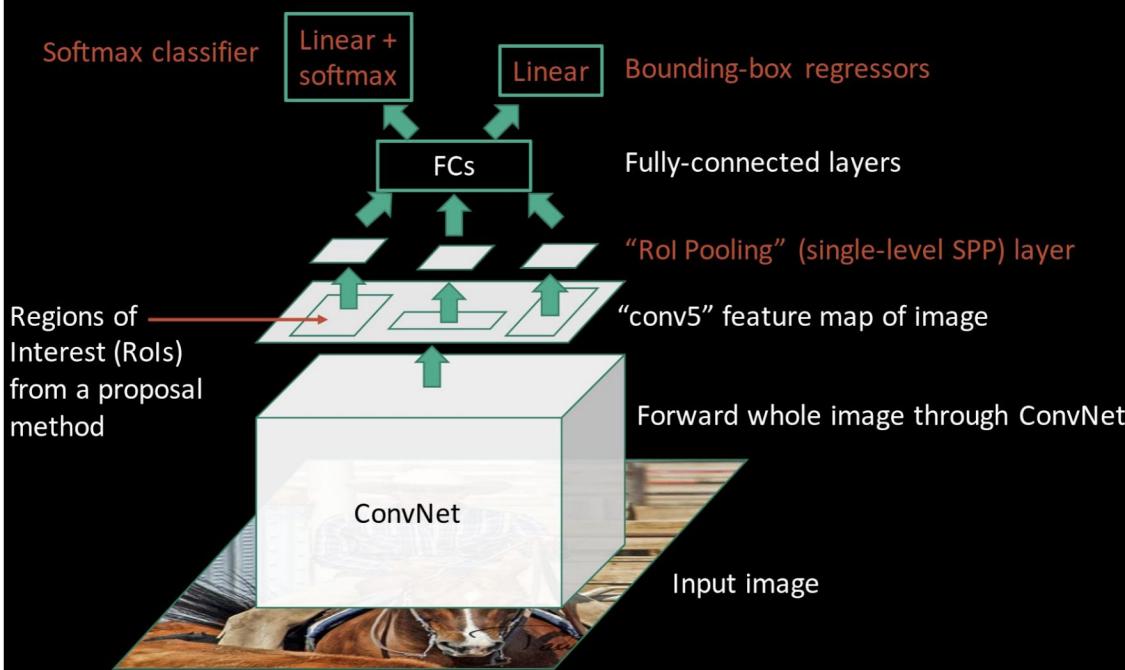


# R-CNN Problems

1. Slow at test-time: need to run full forward pass of CNN for each region proposal
2. SVMs and regressors are post-hoc: CNN features not updated in response to SVMs and regressors
3. Complex multistage training pipeline

# Fast R-CNN

## Fast R-CNN (test time)

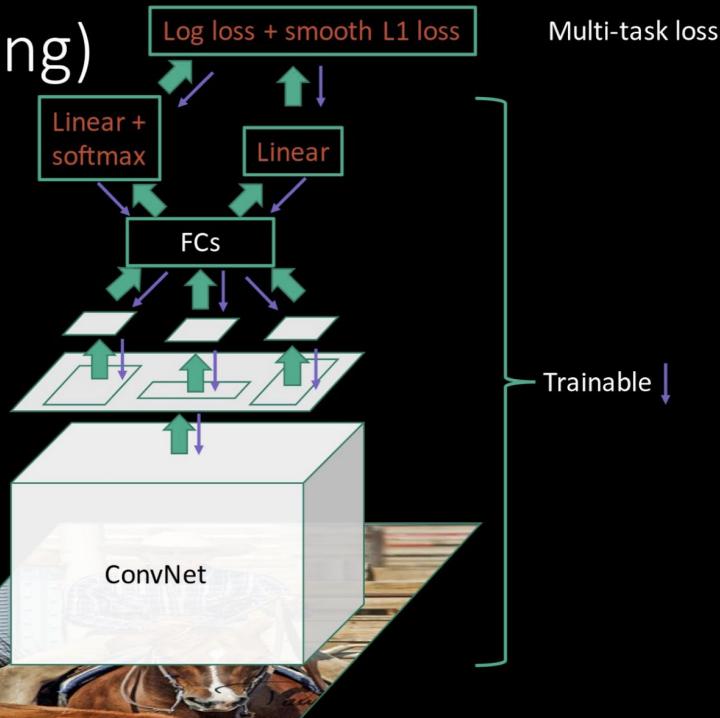


**R-CNN Problem #1:**  
Slow at test-time due to  
independent forward  
passes of the CNN

**Solution:**  
Share computation  
of convolutional  
layers between  
proposals for an  
image

# Fast R-CNN

Fast R-CNN  
(training)



## R-CNN Problem #2:

Post-hoc training: CNN not updated in response to final classifiers and regressors

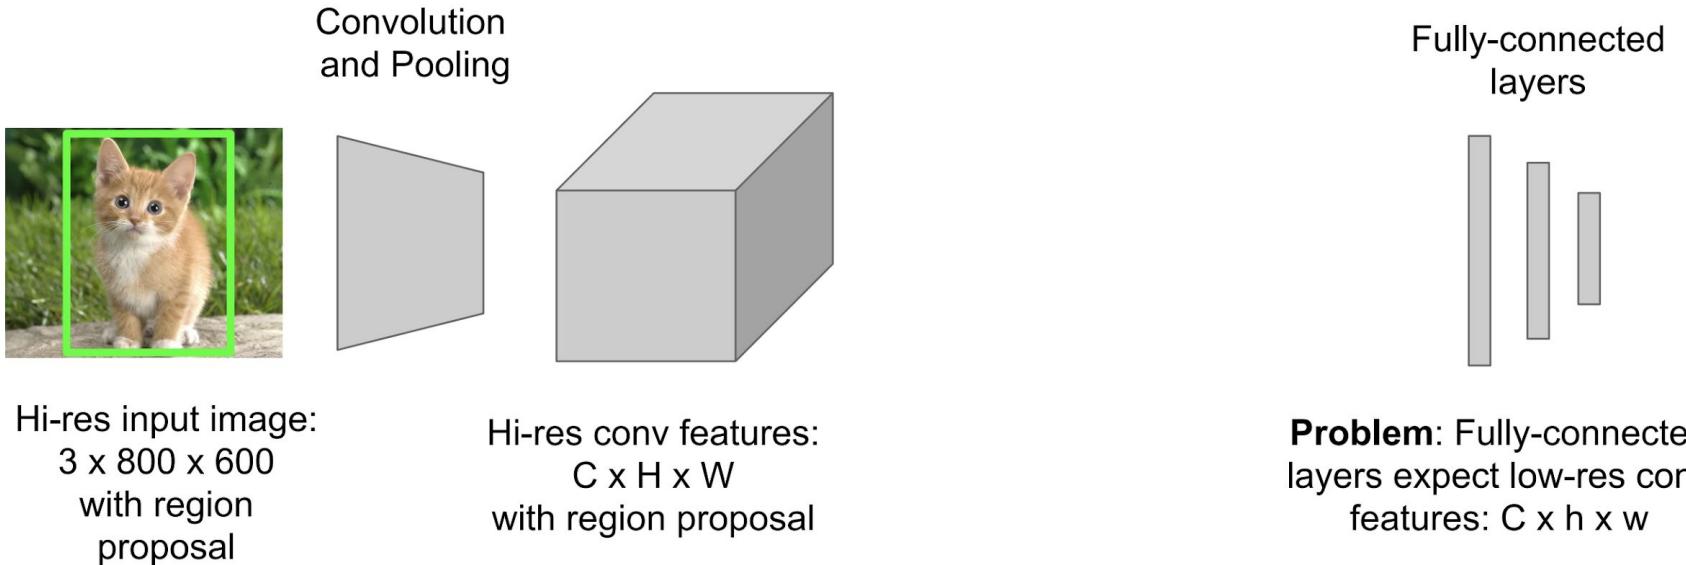
## R-CNN Problem #3:

Complex training pipeline

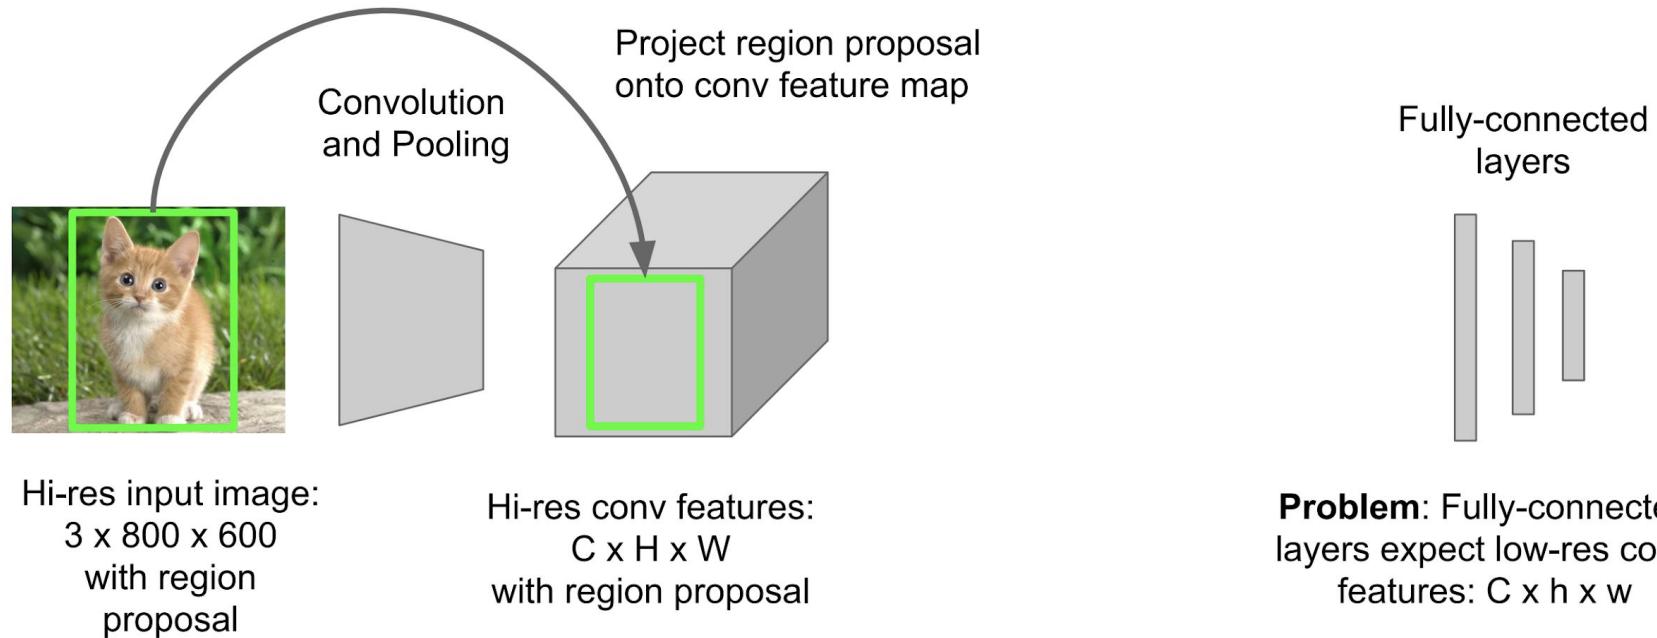
## Solution:

Just train the whole system  
end-to-end all at once!

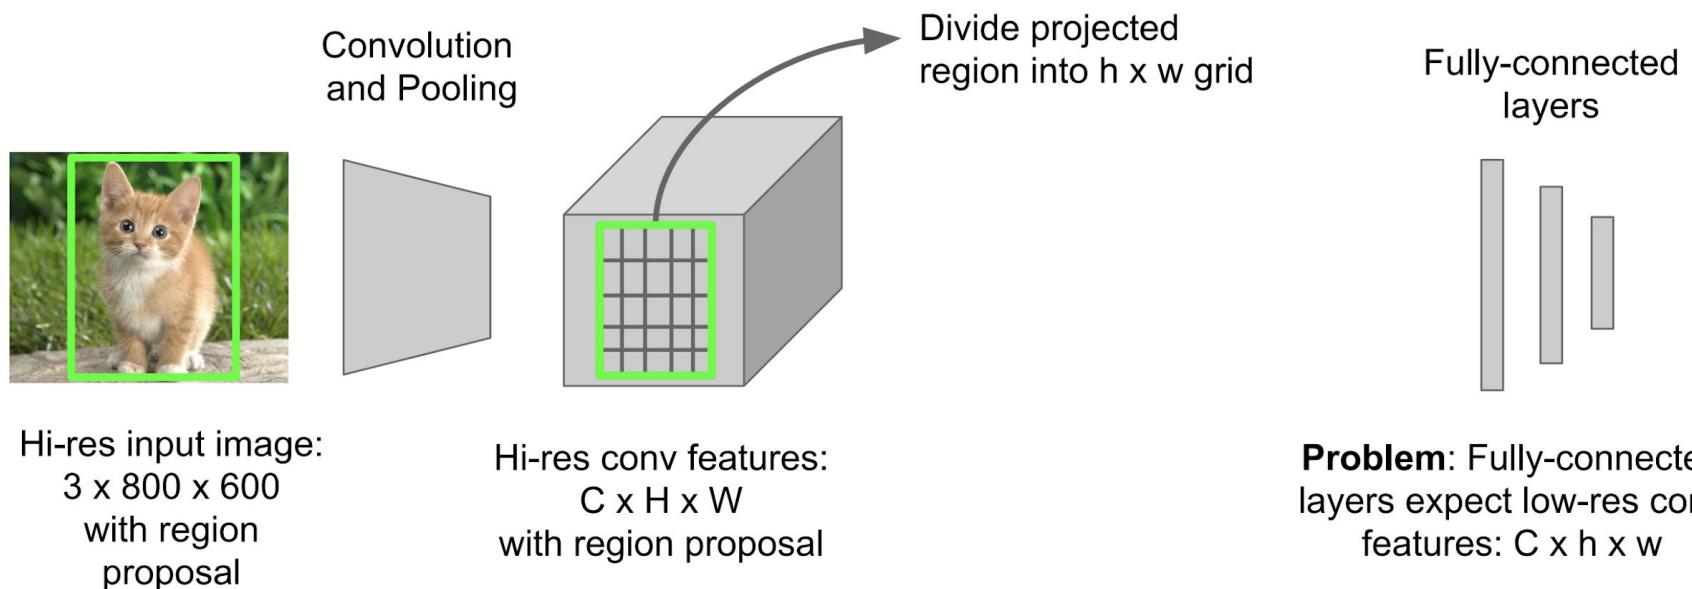
# Fast R-CNN: Region of Interest Pooling



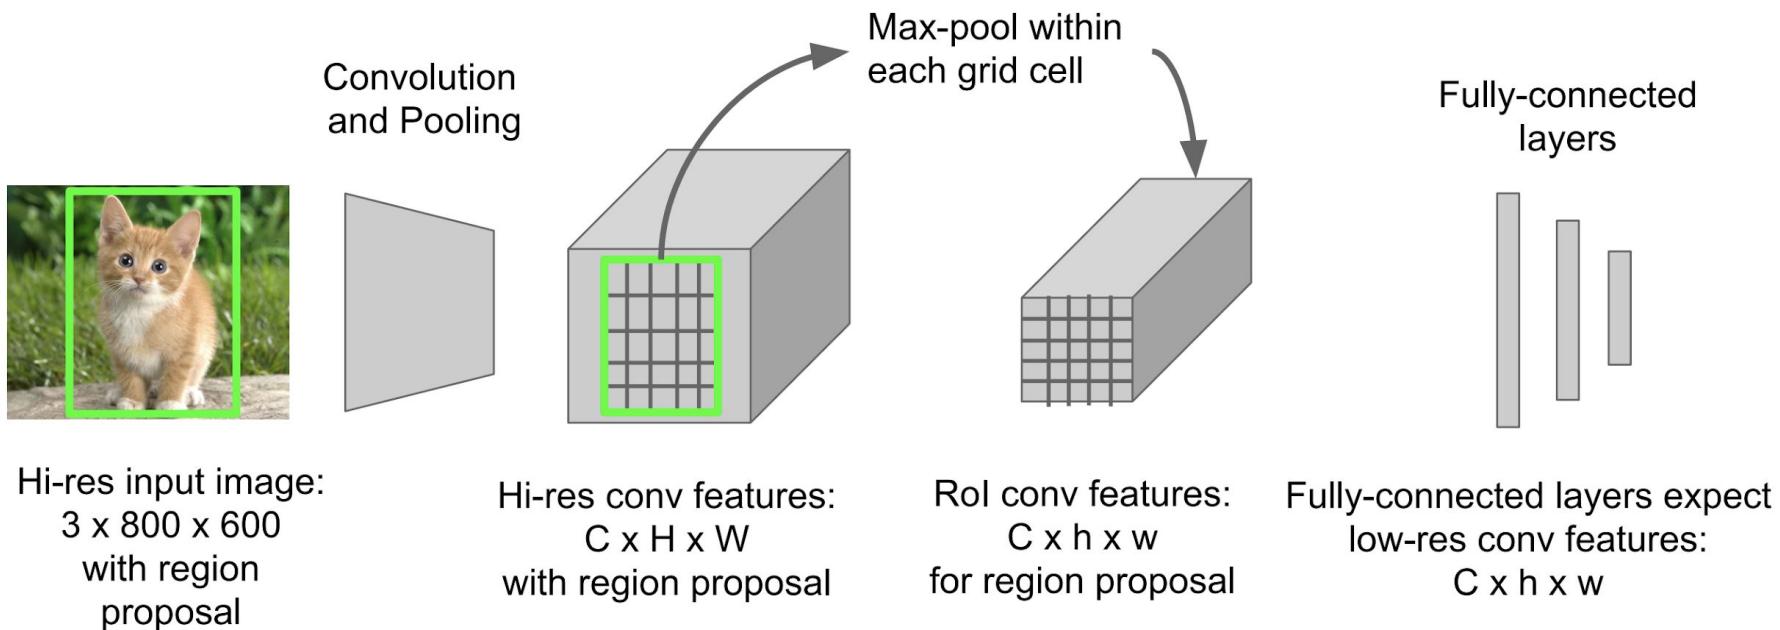
# Fast R-CNN: Region of Interest Pooling



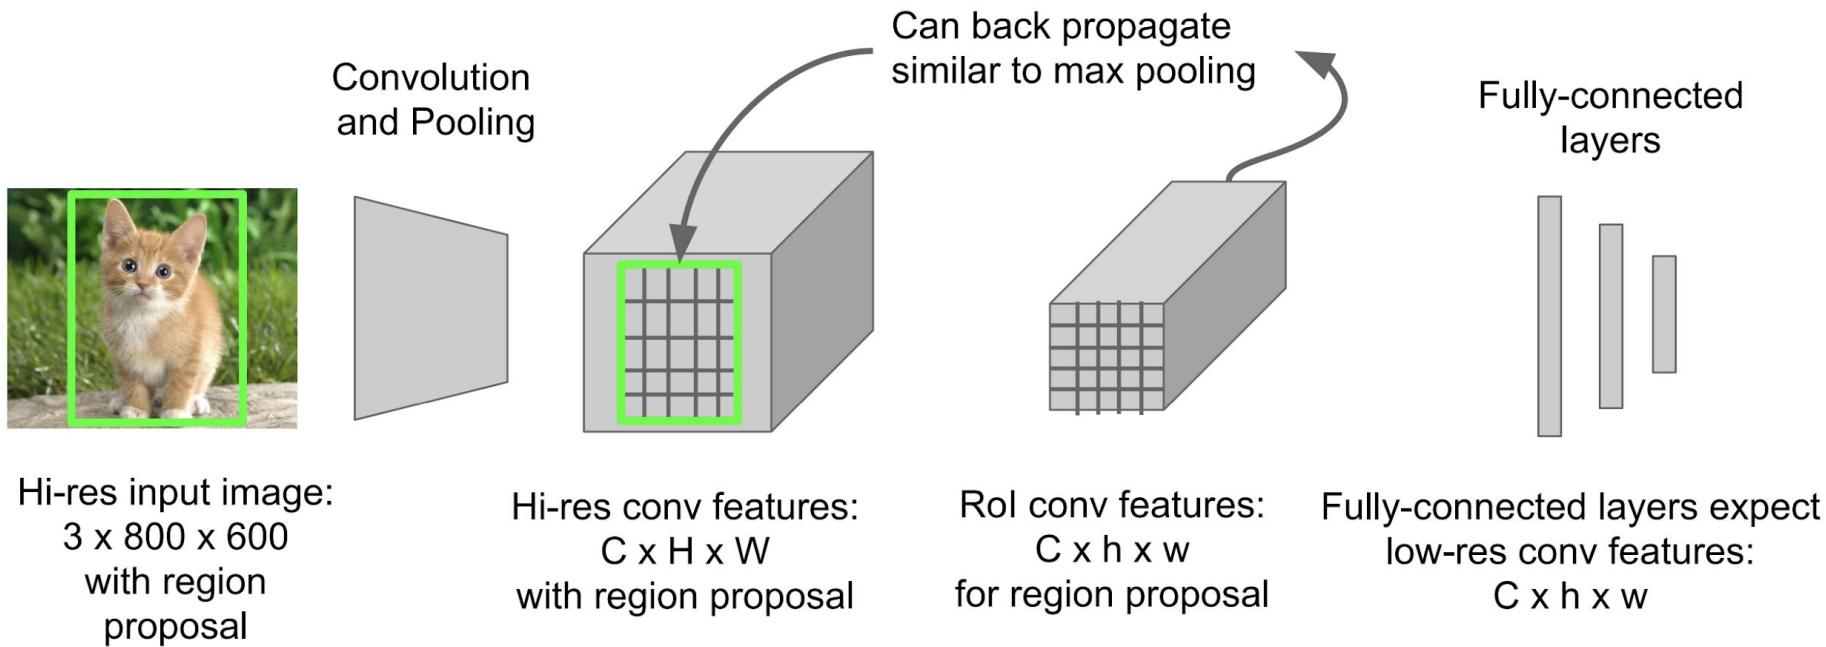
# Fast R-CNN: Region of Interest Pooling



# Fast R-CNN: Region of Interest Pooling



# Fast R-CNN: Region of Interest Pooling



# Fast R-CNN Results

Faster!

FASTER!

Better!

	R-CNN	Fast R-CNN
Training Time:	84 hours	<b>9.5 hours</b>
(Speedup)	1x	<b>8.8x</b>
Test time per image	47 seconds	<b>0.32 seconds</b>
(Speedup)	1x	<b>146x</b>
mAP (VOC 2007)	66.0	<b>66.9</b>

Using VGG-16 CNN on Pascal VOC 2007 dataset

# Fast R-CNN Problem

Test-time speeds don't include region proposals

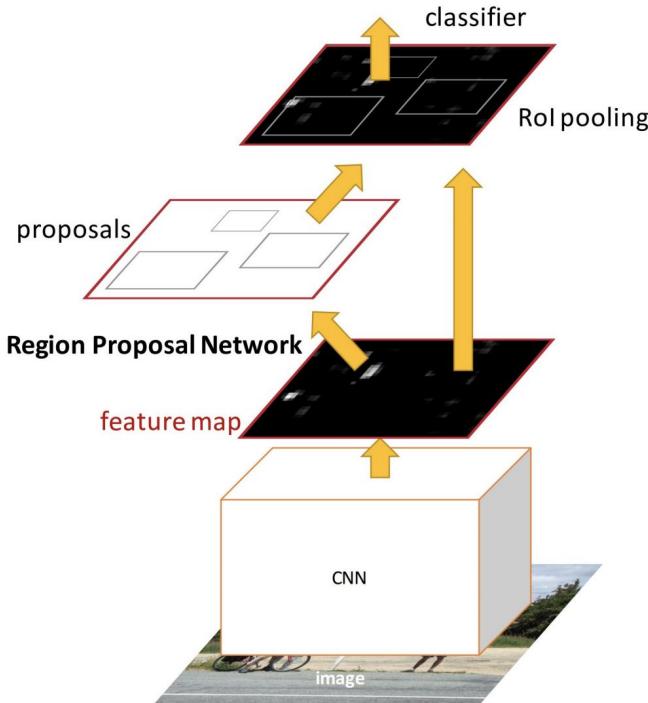
	R-CNN	Fast R-CNN
Test time per image	47 seconds	<b>0.32 seconds</b>
(Speedup)	1x	<b>146x</b>
Test time per image with Selective Search	50 seconds	<b>2 seconds</b>
(Speedup)	1x	<b>25x</b>

# Fast R-CNN Problem Solution

Test-time speeds don't include region proposals  
Just make the CNN do region proposals too!

	R-CNN	Fast R-CNN
Test time per image	47 seconds	<b>0.32 seconds</b>
(Speedup)	1x	<b>146x</b>
Test time per image with Selective Search	50 seconds	<b>2 seconds</b>
(Speedup)	1x	<b>25x</b>

# Faster R-CNN



Insert a **Region Proposal Network (RPN)** after the last convolutional layer

RPN trained to produce region proposals directly; no need for external region proposals!

After RPN, use RoI Pooling and an upstream classifier and bbox regressor just like Fast R-CNN

Ren et al, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015

Slide credit: Ross Girshick

# Faster R-CNN: Region Proposal Network

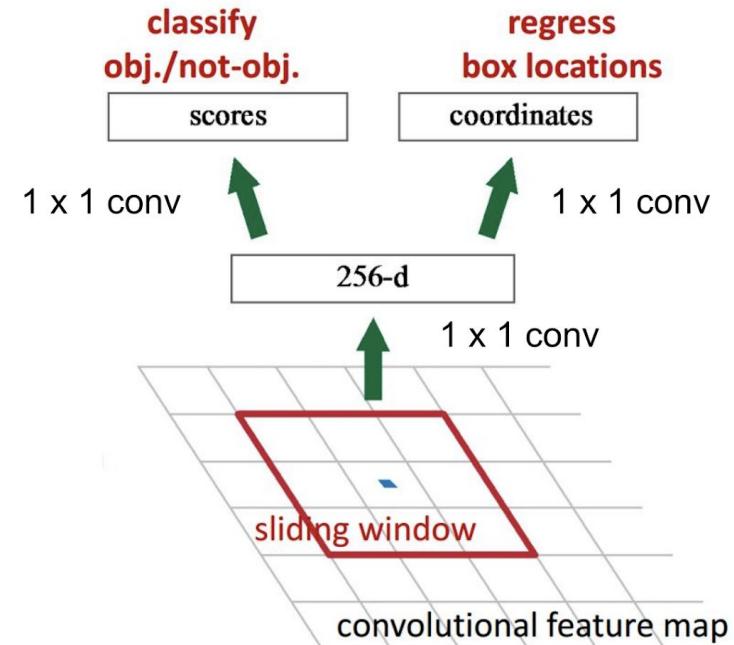
Slide a small window on the feature map

Build a small network for:

- classifying object or not-object, and
- regressing bbox locations

Position of the sliding window provides localization information with reference to the image

Box regression provides finer localization information with reference to this sliding window



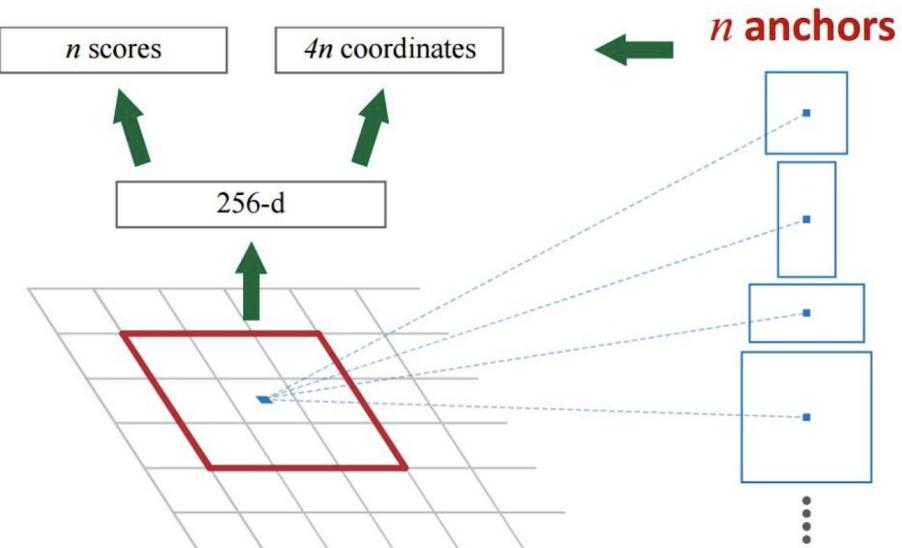
# Faster R-CNN: Region Proposal Network

Use **N anchor boxes** at each location

Anchors are **translation invariant**: use the same ones at every location

Regression gives offsets from anchor boxes

Classification gives the probability that each (regressed) anchor shows an object



# Faster R-CNN: Training

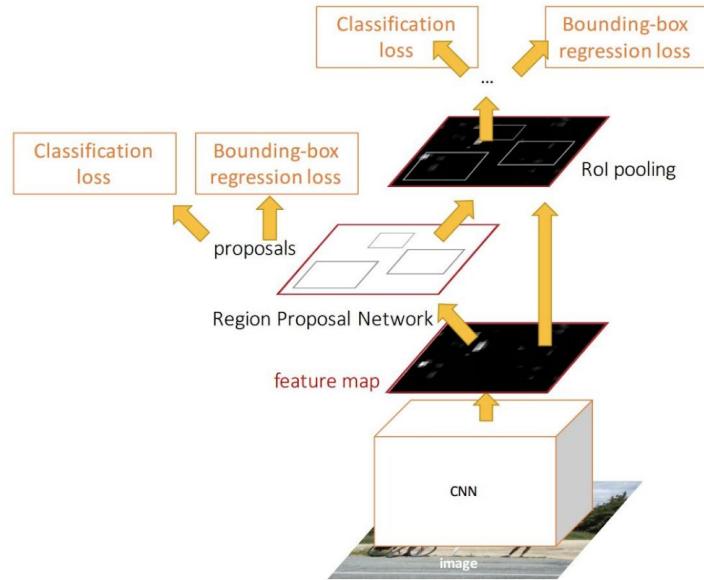
In the paper: Ugly pipeline

- Use alternating optimization to train RPN, then Fast R-CNN with RPN proposals, etc.
- More complex than it has to be

Since publication: Joint training!

One network, four losses

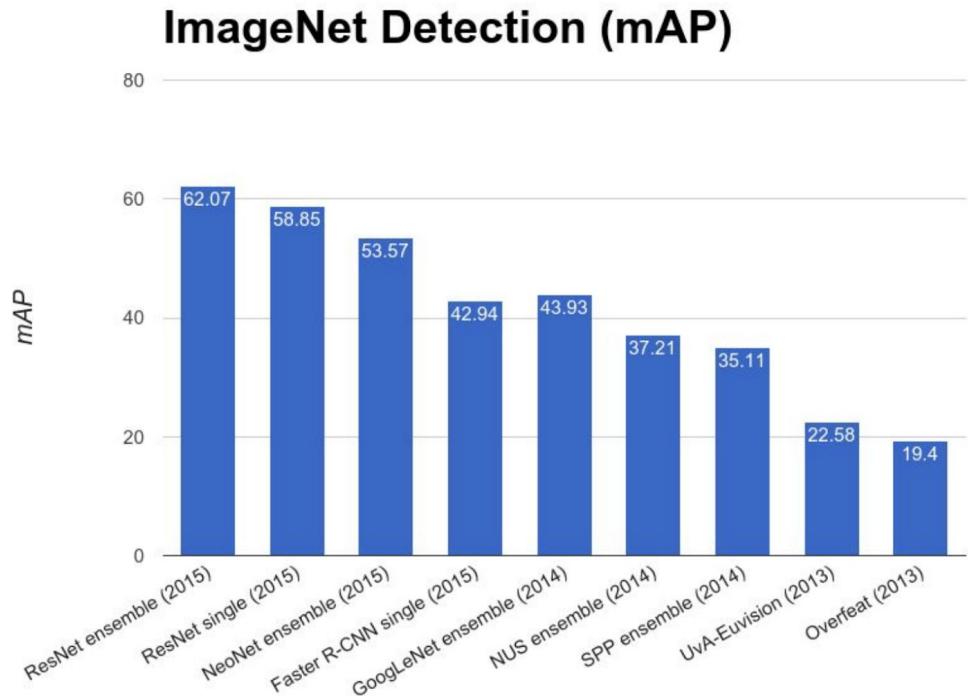
- RPN classification (anchor good / bad)
- RPN regression (anchor  $\rightarrow$  proposal)
- Fast R-CNN classification (over classes)
- Fast R-CNN regression (proposal  $\rightarrow$  box)



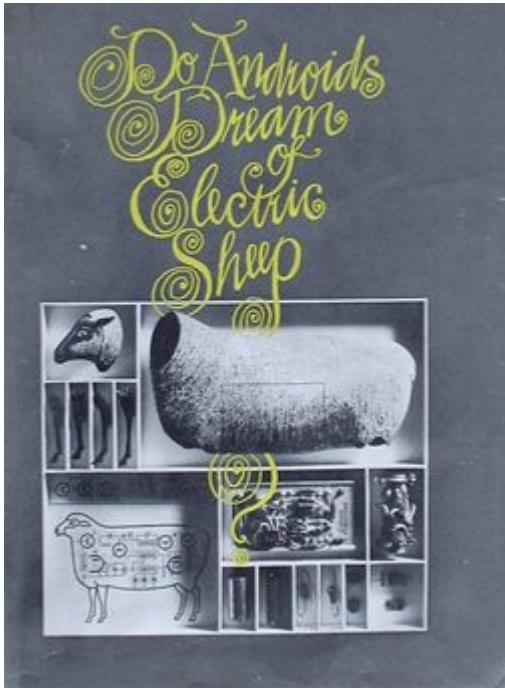
# Faster R-CNN: Results

	<b>R-CNN</b>	<b>Fast R-CNN</b>	<b>Faster R-CNN</b>
Test time per image (with proposals)	50 seconds	2 seconds	<b>0.2 seconds</b>
(Speedup)	1x	25x	<b>250x</b>
mAP (VOC 2007)	66.0	<b>66.9</b>	<b>66.9</b>

# ImageNet Detection 2013 - 2015



# More CV to come...



Next time will:

- try to Only Look Once at an image
- know Do Androids Dream of Electric Sheep?